tensorflow / **tensorflow**

Branch: master ▾   **tensorflow** / tensorflow / examples / udacity / **1_notmnist.ipynb**   Find file   Copy path

zxcqwe4906 Change ndimage.imread to imageio.imread. (#14710)   5fbda9d on 22 Nov 2017

19 contributors

801 lines (800 sloc)   29.5 KB

Deep Learning

Assignment 1

The objective of this assignment is to learn about simple data curation practices, and familiarize you with some of the data we'll be reusing later.

This notebook uses the notMNIST (http://yaroslavvb.blogspot.com/2011/09/notmnist-dataset.html) dataset to be used with python experiments. This dataset is designed to look like the classic MNIST (http://yann.lecun.com/exdb/mnist/) dataset, while looking a little more like real data: it's a harder task, and the data is a lot less 'clean' than MNIST.

```
In [0]:    # These are all the modules we'll be using later. Make sure you can import them
           # before proceeding further.
           from __future__ import print_function
           import imageio
           import matplotlib.pyplot as plt
           import numpy as np
           import os
           import sys
           import tarfile
           from IPython.display import display, Image
           from sklearn.linear_model import LogisticRegression
           from six.moves.urllib.request import urlretrieve
           from six.moves import cPickle as pickle

           # Config the matplotlib backend as plotting inline in IPython
           %matplotlib inline
```

First, we'll download the dataset to our local machine. The data consists of characters rendered in a variety of fonts on a 28x28 image. The labels are limited to 'A' through 'J' (10 classes). The training set has about 500k and the testset 19000 labeled examples. Given these sizes, it should be possible to train models quickly on any machine.

```
In [0]:    url = 'https://commondatastorage.googleapis.com/books1000/'
           last_percent_reported = None
           data_root = '.' # Change me to store data elsewhere

           def download_progress_hook(count, blockSize, totalSize):
             """A hook to report the progress of a download. This is mostly intended for users with
             slow internet connections. Reports every 5% change in download progress.
             """
             global last_percent_reported
             percent = int(count * blockSize * 100 / totalSize)

             if last_percent_reported != percent:
               if percent % 5 == 0:
                 sys.stdout.write("%s%%" % percent)
                 sys.stdout.flush()
               else:
                 sys.stdout.write(".")
                 sys.stdout.flush()

               last_percent_reported = percent

           def maybe_download(filename, expected_bytes, force=False):
             """Download a file if not present, and make sure it's the right size."""
             dest_filename = os.path.join(data_root, filename)
             if force or not os.path.exists(dest_filename):
               print('Attempting to download:', filename)
```

```
      filename, _ = urlretrieve(url + filename, dest_filename, reporthook=download_progress_hook)
      print('\nDownload Complete!')
    statinfo = os.stat(dest_filename)
    if statinfo.st_size == expected_bytes:
      print('Found and verified', dest_filename)
    else:
      raise Exception(
        'Failed to verify ' + dest_filename + '. Can you get to it with a browser?')
    return dest_filename

train_filename = maybe_download('notMNIST_large.tar.gz', 247336696)
test_filename = maybe_download('notMNIST_small.tar.gz', 8458043)
```

```
    Found and verified notMNIST_large.tar.gz
    Found and verified notMNIST_small.tar.gz
```

Extract the dataset from the compressed .tar.gz file. This should give you a set of directories, labeled A through J.

```
In [0]: num_classes = 10
        np.random.seed(133)

        def maybe_extract(filename, force=False):
          root = os.path.splitext(os.path.splitext(filename)[0])[0]  # remove .tar.gz
          if os.path.isdir(root) and not force:
            # You may override by setting force=True.
            print('%s already present - Skipping extraction of %s.' % (root, filename))
          else:
            print('Extracting data for %s. This may take a while. Please wait.' % root)
            tar = tarfile.open(filename)
            sys.stdout.flush()
            tar.extractall(data_root)
            tar.close()
          data_folders = [
            os.path.join(root, d) for d in sorted(os.listdir(root))
            if os.path.isdir(os.path.join(root, d))]
          if len(data_folders) != num_classes:
            raise Exception(
              'Expected %d folders, one per class. Found %d instead.' % (
                num_classes, len(data_folders)))
          print(data_folders)
          return data_folders

        train_folders = maybe_extract(train_filename)
        test_folders = maybe_extract(test_filename)
```

```
['notMNIST_large/A', 'notMNIST_large/B', 'notMNIST_large/C', 'notMNIST_large/D', 'notMNIST_large/E
', 'notMNIST_large/F', 'notMNIST_large/G', 'notMNIST_large/H', 'notMNIST_large/I', 'notMNIST_large
/J']
['notMNIST_small/A', 'notMNIST_small/B', 'notMNIST_small/C', 'notMNIST_small/D', 'notMNIST_small/E
', 'notMNIST_small/F', 'notMNIST_small/G', 'notMNIST_small/H', 'notMNIST_small/I', 'notMNIST_small
/J']
```

Problem 1

Let's take a peek at some of the data to make sure it looks sensible. Each exemplar should be an image of a character A through J rendered in a different font. Display a sample of the images that we just downloaded. Hint: you can use the package IPython.display.

Now let's load the data in a more manageable format. Since, depending on your computer setup you might not be able to fit it all in memory, we'll load each class into a separate dataset, store them on disk and curate them independently. Later we'll merge them into a single dataset of manageable size.

We'll convert the entire dataset into a 3D array (image index, x, y) of floating point values, normalized to have approximately zero mean and standard deviation ~0.5 to make training easier down the road.

A few images might not be readable, we'll just skip them.

```
In [0]: image_size = 28  # Pixel width and height.
        pixel_depth = 255.0  # Number of levels per pixel.

        def load_letter(folder, min_num_images):
          """Load the data for a single letter label."""
```

```python
    image_files = os.listdir(folder)
    dataset = np.ndarray(shape=(len(image_files), image_size, image_size),
                         dtype=np.float32)
    print(folder)
    num_images = 0
    for image in image_files:
      image_file = os.path.join(folder, image)
      try:
        image_data = (imageio.imread(image_file).astype(float) -
                      pixel_depth / 2) / pixel_depth
        if image_data.shape != (image_size, image_size):
          raise Exception('Unexpected image shape: %s' % str(image_data.shape))
        dataset[num_images, :, :] = image_data
        num_images = num_images + 1
      except (IOError, ValueError) as e:
        print('Could not read:', image_file, ':', e, '- it\'s ok, skipping.')

    dataset = dataset[0:num_images, :, :]
    if num_images < min_num_images:
      raise Exception('Many fewer images than expected: %d < %d' %
                      (num_images, min_num_images))

    print('Full dataset tensor:', dataset.shape)
    print('Mean:', np.mean(dataset))
    print('Standard deviation:', np.std(dataset))
    return dataset

def maybe_pickle(data_folders, min_num_images_per_class, force=False):
  dataset_names = []
  for folder in data_folders:
    set_filename = folder + '.pickle'
    dataset_names.append(set_filename)
    if os.path.exists(set_filename) and not force:
      # You may override by setting force=True.
      print('%s already present - Skipping pickling.' % set_filename)
    else:
      print('Pickling %s.' % set_filename)
      dataset = load_letter(folder, min_num_images_per_class)
      try:
        with open(set_filename, 'wb') as f:
          pickle.dump(dataset, f, pickle.HIGHEST_PROTOCOL)
      except Exception as e:
        print('Unable to save data to', set_filename, ':', e)

  return dataset_names

train_datasets = maybe_pickle(train_folders, 45000)
test_datasets = maybe_pickle(test_folders, 1800)
```

```
notMNIST_large/A
Could not read: notMNIST_large/A/Um9tYW5hIEJvbGQucGZi.png : cannot identify image file - it's ok,
skipping.
Could not read: notMNIST_large/A/RnJlaWdodERpc3BCb29rSXRhbGljLnR0Zg==.png : cannot identify image
file - it's ok, skipping.
Could not read: notMNIST_large/A/SG90IE1lc3RhcmQgQlROIFBvc3Rlci50dGY=.png : cannot identify image
file - it's ok, skipping.
Full dataset tensor: (52909, 28, 28)
Mean: -0.12848
Standard deviation: 0.425576
notMNIST_large/B
Could not read: notMNIST_large/B/TmlraXNFRi1lTZW1pQm9sZEl0YWxpYy5vdGY=.png : cannot identify image
file - it's ok, skipping.
Full dataset tensor: (52911, 28, 28)
Mean: -0.00755947
Standard deviation: 0.417272
notMNIST_large/C
Full dataset tensor: (52912, 28, 28)
Mean: -0.142321
Standard deviation: 0.421305
notMNIST_large/D
Could not read: notMNIST_large/D/VHJhbnNpdCBCb2xkLnR0Zg==.png : cannot identify image file - it's
ok, skipping.
Full dataset tensor: (52911, 28, 28)
Mean: -0.0574553
Standard deviation: 0.434072
notMNIST_large/E
Full dataset tensor: (52912, 28, 28)
Mean: -0.0701406
Standard deviation: 0.42882
```

```
notMNIST_large/F
Full dataset tensor: (52912, 28, 28)
Mean: -0.125914
Standard deviation: 0.429645
notMNIST_large/G
Full dataset tensor: (52912, 28, 28)
Mean: -0.0947771
Standard deviation: 0.421674
notMNIST_large/H
Full dataset tensor: (52912, 28, 28)
Mean: -0.0687667
Standard deviation: 0.430344
notMNIST_large/I
Full dataset tensor: (52912, 28, 28)
Mean: 0.0307405
Standard deviation: 0.449686
notMNIST_large/J
Full dataset tensor: (52911, 28, 28)
Mean: -0.153479
Standard deviation: 0.397169
notMNIST_small/A
Could not read: notMNIST_small/A/RGVtb2NyYXRpY2FCb2xkT2xkc3R5bGUgQm9sZC50dGY=.png : cannot identif
y image file - it's ok, skipping.
Full dataset tensor: (1872, 28, 28)
Mean: -0.132588
Standard deviation: 0.445923
notMNIST_small/B
Full dataset tensor: (1873, 28, 28)
Mean: 0.00535619
Standard deviation: 0.457054
notMNIST_small/C
Full dataset tensor: (1873, 28, 28)
Mean: -0.141489
Standard deviation: 0.441056
notMNIST_small/D
Full dataset tensor: (1873, 28, 28)
Mean: -0.0492094
Standard deviation: 0.460477
notMNIST_small/E
Full dataset tensor: (1873, 28, 28)
Mean: -0.0598952
Standard deviation: 0.456146
notMNIST_small/F
Could not read: notMNIST_small/F/Q3Jvc3NvdmVyIEJvbGRPYmxpcXVlLnR0Zg==.png : cannot identify image
file - it's ok, skipping.
Full dataset tensor: (1872, 28, 28)
Mean: -0.118148
Standard deviation: 0.451134
notMNIST_small/G
Full dataset tensor: (1872, 28, 28)
Mean: -0.092519
Standard deviation: 0.448468
notMNIST_small/H
Full dataset tensor: (1872, 28, 28)
Mean: -0.0586729
Standard deviation: 0.457387
notMNIST_small/I
Full dataset tensor: (1872, 28, 28)
Mean: 0.0526481
Standard deviation: 0.472657
notMNIST_small/J
Full dataset tensor: (1872, 28, 28)
Mean: -0.15167
Standard deviation: 0.449521
```

Problem 2

Let's verify that the data still looks good. Displaying a sample of the labels and images from the ndarray. Hint: you can use matplotlib.pyplot.

Problem 3

Another check: we expect the data to be balanced across classes. Verify that.

Merge and prune the training data as needed. Depending on your computer setup, you might not be able to fit it all in memory, and you can tune `train_size` as needed. The labels will be stored into a separate array of integers 0 through 9.

Also create a validation dataset for hyperparameter tuning.

```python
In [0]: def make_arrays(nb_rows, img_size):
          if nb_rows:
            dataset = np.ndarray((nb_rows, img_size, img_size), dtype=np.float32)
            labels = np.ndarray(nb_rows, dtype=np.int32)
          else:
            dataset, labels = None, None
          return dataset, labels

        def merge_datasets(pickle_files, train_size, valid_size=0):
          num_classes = len(pickle_files)
          valid_dataset, valid_labels = make_arrays(valid_size, image_size)
          train_dataset, train_labels = make_arrays(train_size, image_size)
          vsize_per_class = valid_size // num_classes
          tsize_per_class = train_size // num_classes

          start_v, start_t = 0, 0
          end_v, end_t = vsize_per_class, tsize_per_class
          end_l = vsize_per_class+tsize_per_class
          for label, pickle_file in enumerate(pickle_files):
            try:
              with open(pickle_file, 'rb') as f:
                letter_set = pickle.load(f)
                # let's shuffle the letters to have random validation and training set
                np.random.shuffle(letter_set)
                if valid_dataset is not None:
                  valid_letter = letter_set[:vsize_per_class, :, :]
                  valid_dataset[start_v:end_v, :, :] = valid_letter
                  valid_labels[start_v:end_v] = label
                  start_v += vsize_per_class
                  end_v += vsize_per_class

                train_letter = letter_set[vsize_per_class:end_l, :, :]
                train_dataset[start_t:end_t, :, :] = train_letter
                train_labels[start_t:end_t] = label
                start_t += tsize_per_class
                end_t += tsize_per_class
            except Exception as e:
              print('Unable to process data from', pickle_file, ':', e)
              raise

          return valid_dataset, valid_labels, train_dataset, train_labels


        train_size = 200000
        valid_size = 10000
        test_size = 10000

        valid_dataset, valid_labels, train_dataset, train_labels = merge_datasets(
          train_datasets, train_size, valid_size)
        _, _, test_dataset, test_labels = merge_datasets(test_datasets, test_size)

        print('Training:', train_dataset.shape, train_labels.shape)
        print('Validation:', valid_dataset.shape, valid_labels.shape)
        print('Testing:', test_dataset.shape, test_labels.shape)
```

```
Training (200000, 28, 28) (200000,)
Validation (10000, 28, 28) (10000,)
Testing (10000, 28, 28) (10000,)
```

Next, we'll randomize the data. It's important to have the labels well shuffled for the training and test distributions to match.

```python
In [0]: def randomize(dataset, labels):
          permutation = np.random.permutation(labels.shape[0])
          shuffled_dataset = dataset[permutation,:,:]
          shuffled_labels = labels[permutation]
```

```
    return shuffled_dataset, shuffled_labels
train_dataset, train_labels = randomize(train_dataset, train_labels)
test_dataset, test_labels = randomize(test_dataset, test_labels)
valid_dataset, valid_labels = randomize(valid_dataset, valid_labels)
```

Problem 4

Convince yourself that the data is still good after shuffling!

Finally, let's save the data for later reuse:

```
In [0]: pickle_file = os.path.join(data_root, 'notMNIST.pickle')

        try:
          f = open(pickle_file, 'wb')
          save = {
            'train_dataset': train_dataset,
            'train_labels': train_labels,
            'valid_dataset': valid_dataset,
            'valid_labels': valid_labels,
            'test_dataset': test_dataset,
            'test_labels': test_labels,
            }
          pickle.dump(save, f, pickle.HIGHEST_PROTOCOL)
          f.close()
        except Exception as e:
          print('Unable to save data to', pickle_file, ':', e)
          raise
```

```
In [0]: statinfo = os.stat(pickle_file)
        print('Compressed pickle size:', statinfo.st_size)
```

```
Compressed pickle size: 718193801
```

Problem 5

By construction, this dataset might contain a lot of overlapping samples, including training data that's also contained in the validation and test set! Overlap between training and test can skew the results if you expect to use your model in an environment where there is never an overlap, but are actually ok if you expect to see training samples recur when you use it. Measure how much overlap there is between training, validation and test samples.

Optional questions:

- What about near duplicates between datasets? (images that are almost identical)
- Create a sanitized validation and test set, and compare your accuracy on those in subsequent assignments.

Problem 6

Let's get an idea of what an off-the-shelf classifier can give you on this data. It's always good to check that there is something to learn, and that it's a problem that is not so trivial that a canned solution solves it.